CS547 HW3

Bo Chen

September 26 2019

The test accuracy is 81.0%, which is achieved after 146 epochs. The input images are flipped and cropped for data augmentation. The implementation consists of 5 convolution layers followed by 3 fully connected layers. The kernel size, the stride length, and the padding of each convolution layer is $(3, 3)$, 1 and 0. The number of channels of convolution layers are 16, 32, 64, 128, and 128. Batch normalization is applied to the output of each convolution layer and the dropout layer of 0.5 is applied to the first, the third and the fifth output of batch normalization. The size of the fully connected layers are 1152×128, 128×64, and 64×10. The adam optimizer is used for optimization and the learning rate is set to 0.001.